

Asking AI Better Questions

A practical guide for business analysis,
decision support, and responsible implementation

AI should not merely produce an answer.

It should help design a decision process that knows when to trust, verify, question, escalate, and stop.

Contents

1	The right mental model	2
1.1	Make room for uncertainty	2
2	How to use this guide	3
3	The ten-question core checklist	3
4	Strategy and problem definition	3
5	Context, data, and evidence	4
6	Accuracy, assumptions, and anti-sycophancy	5
7	Risk, bias, and ethics	5
8	Legal, intellectual property, privacy, and compliance	6
9	Evaluation, testing, and decision support	6
10	Operations, implementation, and AI agents	7
11	Cost, ROI, and vendor realism	8
12	People, process, and change	8
13	Quality, alternatives, and communication	9
14	Reproducibility and organisational memory	9
15	Better prompting: ask the AI how to use it	10
16	Reusable prompt templates	10
16.1	The default business prompt	10
16.2	The anti-sycophancy prompt	11
16.3	The pilot-design prompt	11
16.4	The agent-permissions prompt	11
16.5	The source-grounded prompt	11
17	Before acting: the final gate	11

1 The right mental model

Treat AI as an extremely fast, articulate junior analyst: useful for drafting, comparison, pattern-finding, summarisation, coding, and generating alternatives—but dependent on context, unreliable at the edges, persuasive when wrong, and unable to own consequences.

The model may produce a polished answer even when the problem is poorly defined, the data is incomplete, the sources are stale, or the decision requires authority that the model does not possess. Its output is **work product, not evidence**.

The quality of an AI-supported decision depends on more than the prompt. It depends on:

- the business problem and the decision being improved;
- the context, data, sources, tools, and permissions available;
- the model's ability to represent uncertainty and recognise its limits;
- the verification, evaluation, and monitoring process;
- the workflow, incentives, people, and accountability surrounding the system.

Central rule. Ask AI for analysis, alternatives, tests, drafts, and structured reasoning. Do not outsource the ownership of a material decision merely because the model can express one confidently.

1.1 Make room for uncertainty

AI systems are generally optimised to be useful, fluent, and responsive. In practice, this can create **sycophancy**: the system mirrors the user's framing, accepts questionable premises, agrees too readily, and fills gaps with a plausible narrative. A user who asks for one confident answer may receive one even when several interpretations are possible.

This is not solved by asking the model to “be confident” or by asking it to provide a numerical confidence score. Self-reported confidence is not automatically calibrated. A better approach is to make uncertainty operational and visible.

Do not force completion. If the evidence is insufficient, the correct output may be a question, a placeholder, a range of possibilities, an explicit “unknown”, or a recommendation to escalate. A visibly incomplete answer is safer than a fabricated complete one.

Useful instructions include:

- “Challenge my framing before answering.”
- “Give me several plausible interpretations, not one convenient story.”
- “Separate facts, sourced claims, assumptions, inferences, and unknowns.”
- “Use **[VERIFY]**, **[MISSING DATA]**, **[ASSUMPTION]**, and **[SOURCE NEEDED]** instead of inventing details.”
- “State what evidence would change your conclusion.”
- “If the information is insufficient, ask for clarification or recommend escalation.”
- “Do not invent citations, quotations, calculations, legal rules, customer facts, or experimental results.”

A good default attitude: be open to the model being useful, but make it earn trust claim by claim.

2 How to use this guide

The questions below are a reusable prompt library. They are not a script that should be pasted mechanically into every conversation.

1. Start with the ten-question core checklist.
2. Add only the sections relevant to the decision and its risk.
3. Ask no more than five clarifying questions at a time, ranked by expected impact. If they remain unanswered, proceed using explicit assumptions.
4. Require stronger evidence when the output will be published, priced, signed, regulated, sent to a customer, used in hiring, or automated.
5. Use deterministic tools—formulas, calculators, databases, code, lookup tables, and policy engines—when they are safer than language generation.
6. Keep a human accountable for money, law, employment, safety, medical decisions, reputation, and irreversible actions.

Suggested answer structure for AI: Decision needed; known facts; missing information; assumptions; alternatives; evidence; risks; human approvals; test plan; recommendation or next action.

3 The ten-question core checklist

Use these questions as a compact starting point for almost any business task:

1. What decision or process are we trying to improve?
2. What would change if the answer were correct?
3. What information is missing, and which missing item matters most?
4. Which claims are facts, sourced evidence, assumptions, inferences, or guesses?
5. What are the strongest alternative interpretations or options?
6. What could go wrong, and who would be affected?
7. Which parts should use a calculator, database, code, search engine, or policy rule instead of you?
8. What must a human approve, and who owns the outcome?
9. What is the cheapest test that could disprove this recommendation within one or two weeks?
10. What metric, owner, review date, and stopping rule determine whether we continue?

4 Strategy and problem definition

Before discussing models, clarify whether there is a valuable problem at all.

- What problem am I actually trying to solve, and is AI the right tool versus a spreadsheet, an API, a rule, or an existing process?
- Restate my goal in one sentence, then list the assumptions I am smuggling in.
- What decision, behaviour, or workflow should change if this initiative succeeds?
- Who makes that decision today, who will use the output, and who bears the consequences if it is wrong?
- What would have to be true for this idea or plan to work?
- What is the cost of doing nothing, and what is the cost of doing the wrong thing?
- Argue against this plan. Give the five strongest reasons a sceptical investor, board, customer, employee, or regulator would reject it.
- Where does this accelerate our current strategy, and where might it force a change in the business model?
- Are we automating an existing process that should instead be redesigned or removed?
- What is the decision error we are trying to reduce: delay, inconsistency, missed opportunities, false alarms, fraud, waste, or something else?
- What would a successful outcome look like in operational terms, not merely in model metrics?
- Compare this initiative with doing nothing and with improving the current manual process.

5 Context, data, and evidence

Language models are highly sensitive to what is placed in their context. A better question cannot compensate for absent or inaccessible evidence.

- What information is missing from what I gave you that would materially change the answer? Rank the missing information by value.
- Which claims are directly supported by my inputs, which require external sources, and which are your inferences?
- What sources, data types, internal documents, or subject-matter experts would make this more reliable?
- For every important claim, give the supporting source, publication date, jurisdiction, and verification status.
- Is this answer predictable from historical patterns, or is this a rare, novel, or distribution-shifted situation where AI may be weak?
- What would “good data” look like for this use case, and what quality problems could silently ruin it?
- If my CRM, files, and notes are messy or incomplete, how should I distinguish absence of evidence from evidence of absence?
- What sampling bias, missing-not-at-random data, label error, or historical bias may be present?
- Which internal documents are authoritative, and which may be obsolete, contradictory, or merely opinion?
- What information should be retrieved from a database or document system rather than generated from general model knowledge?

- What date range, language, geography, customer segment, or operating condition does the evidence actually cover?
- What would an evidence table look like with columns for claim, source, date, strength, uncertainty, and open question?

6 Accuracy, assumptions, and anti-sycophancy

The purpose of these prompts is not to make the model sound cautious. It is to expose the structure of the uncertainty.

- Separate every important statement into fact, sourced claim, assumption, inference, recommendation, or unknown.
- What are the top ways this answer could be wrong or misleading?
- What alternative interpretations am I not seeing?
- What is the strongest argument against your current conclusion?
- What evidence, assumption, or boundary condition would make you change your conclusion?
- Give me three plausible explanations and the evidence that would distinguish among them.
- Give me several options with trade-offs rather than one answer. Do not collapse uncertainty into a single recommendation unless the decision rule is explicit.
- Which parts of this answer are genuinely uncertain, and what should remain as a placeholder until verified?
- If two interpretations are both plausible, show both and explain what information would resolve the ambiguity.
- Try to disagree with me. Which premise would you challenge first?
- Check whether you are simply accepting my wording, goals, or implied conclusion. Rewrite the problem neutrally.
- Do not fill gaps with plausible details. Mark them as **[TO VERIFY]** or **[UNKNOWN]**.
- Do not invent sources, quotations, case studies, numbers, names, or legal requirements. If you cannot verify one, say so.

Better than “rate your confidence”:

“Separate facts, assumptions, inferences, and unknowns. For each important claim, explain the evidence supporting it, the main uncertainty, and what would change your conclusion. Use qualitative support levels unless a numerical probability is grounded in data or a calibrated model.”

7 Risk, bias, and ethics

Risk analysis should include not only whether the model is accurate, but also who gains power, who loses access, and who must absorb the errors.

- Who could be harmed if this is wrong, incomplete, or biased?
- Which groups, regions, languages, customer types, or atypical cases might this systematically underserve or misrepresent?
- What hidden incentives or default assumptions are baked into this recommendation?

- What should never be automated here, and where must a human decide?
- What are the important false positives and false negatives? Which is more costly?
- What failure modes become more damaging when the system is used at scale?
- Could automation bias cause people to accept the output too readily?
- Could the system create surveillance, deskilling, discrimination, exclusion, or a transfer of risk to less powerful workers?
- If this output leaked or was screenshotted, what reputational, contractual, or safety risk would it create?
- How can an affected customer, employee, or citizen challenge or appeal the outcome?
- Are people entitled to know that AI is involved, and what explanation should they receive?
- Whose values are absent from the problem formulation?

8 Legal, intellectual property, privacy, and compliance

Legal questions are jurisdiction-specific and time-sensitive. AI can help identify issues and prepare questions for a qualified professional; it should not be treated as the final legal authority.

- Which jurisdiction, date, sector, and risk category apply to this use case?
- Given GDPR, the EU AI Act, employment law, financial regulation, medical rules, advertising rules, or other applicable requirements, what must be checked before use?
- Is a data-protection impact assessment, risk assessment, human-rights assessment, or sector approval required?
- Who is the controller, processor, deployer, provider, or responsible decision-maker?
- What personal, health, financial, confidential, or commercially sensitive data should not be pasted into a public AI tool?
- What redaction, access-control, retention, deletion, and cross-border-transfer rules should apply?
- Does this output contain material that should not be treated as original or as our intellectual property?
- What copyright, trademark, publicity, licensing, or attribution issues should be checked before publishing or selling this?
- Who should own the output contractually if a vendor, employee, contractor, or model provider is involved?
- What contractual promises does the vendor make about training data, confidentiality, retention, uptime, indemnity, and model changes?
- What audit trail would a regulator, customer, court, or internal reviewer reasonably expect?
- What evidence is needed before making a legal claim rather than merely identifying a legal question?

9 Evaluation, testing, and decision support

This is the section that turns a plausible answer into a testable intervention.

- What is the baseline and counterfactual: what happens today, and what would happen without this intervention?
- Define a test set containing normal cases, edge cases, adversarial cases, and known failures.
- What acceptance criteria must be met before deployment?
- Which errors matter most to the business, and how will they be measured?
- What is the cheapest test that could disprove this recommendation in one or two weeks?
- What result would justify scaling, redesigning, pausing, or stopping?
- When should the system abstain, ask for clarification, or escalate to a human?
- How will we measure whether the model improves the actual decision rather than merely its own benchmark?
- What user adoption, override, correction, exception, and abandonment rates should we monitor?
- What would indicate data drift, concept drift, calibration failure, or a change in the underlying business process?
- What metric should move in 30, 60, and 90 days if this is working?
- Can the result be checked using a second model, a deterministic calculation, a sample audit, or an independent expert?
- How will we distinguish real value from temporary novelty, unpaid review work, or a one-off workaround?
- What incident threshold triggers investigation, rollback, retraining, or suspension?

Evaluation principle. A high model score is not evidence of business value. The relevant question is whether the intervention changes a decision or process in a measurable, safe, and economically worthwhile way.

10 Operations, implementation, and AI agents

An AI initiative becomes operational when its output enters a workflow. This is where permissions, exceptions, failures, and human responsibility must be designed.

- Break this into a workflow: what should AI draft, what must a human approve, and what must remain off-limits?
- What permissions, systems, and data access would an AI agent need?
- What should the agent never be allowed to read, write, send, approve, purchase, delete, or publish?
- Can permissions be limited according to least privilege and granted only when needed?
- How should the system handle vague, incomplete, contradictory, or conflicting user requests?
- What happens when the model is down, stale, unavailable, or confidently wrong mid-process?
- What manual fallback keeps the business operating during an outage?
- How are exceptions routed, reviewed, resolved, and fed back into the system?
- What monitoring would tell us that the quality is drifting or users are quietly routing around

the tool?

- If we deploy this, what breaks in adjacent teams, systems, incentives, or customer journeys?
- Could an email, webpage, document, or retrieved passage contain a prompt injection? How will untrusted content be separated from instructions?
- Can every agent action be logged, reconstructed, reversed, and safely repeated?
- What approval gate is required before an irreversible or high-impact action?
- Who owns the exception process when the model and the human disagree?

11 Cost, ROI, and vendor realism

AI projects often look inexpensive because the pilot budget omits the work required to make the system reliable.

- What is the full cost, including data cleanup, integration, security, review time, training, monitoring, error handling, and customer recovery?
- What is the fully loaded cost per useful decision, not merely the subscription or API price?
- Where will the system create more work—review, exception handling, escalation, correction, or communication—even if it saves drafting time?
- Apply a conservative haircut to these benefit claims. What remains under pessimistic assumptions?
- Compare the economics with the current manual process and with a simpler deterministic alternative.
- Will the workflow we are optimising still exist in three years, or are we automating something that should be redesigned?
- Where does vendor lock-in, model concentration, price volatility, or service-dependency risk appear?
- What happens if the vendor changes the model, raises prices, removes a feature, suffers an outage, or exits the market?
- What portability, export, fallback, and termination rights do we need?
- What would make this initiative not worth the money? Define the stopping condition in advance.

12 People, process, and change

Adoption is not a communications problem added after the technology is built. It is part of the system design.

- What skills does my team still need if AI produces the first draft or recommendation?
- Where will this create additional work for employees, even if it saves time elsewhere?
- Who performs the hidden review, correction, labelling, escalation, and customer-recovery work?
- What improves for the people doing the work? State the employee deal explicitly.
- How might employees interpret this tool—as assistance, surveillance, deskilling, or a head-count plan?

- Who owns the outcome if the AI is wrong, and do they have the authority and resources to intervene?
- How will users challenge, correct, and learn from the system without being punished for reporting failures?
- How do we keep institutional knowledge from rotting if people stop writing or reasoning through the first version themselves?
- Which skills, judgement, and professional roles might decay through over-reliance?
- What training, incentives, time, and local process changes are needed for adoption?
- What should the internal AI-use policy say about this task?
- How will feedback from frontline workers change the workflow, data, model, or governance?
- Does the initiative redistribute power or visibility in the organisation? Who gains and who loses discretion?

13 Quality, alternatives, and communication

Use AI to improve the quality of the decision process, not merely to produce more text.

- Give me three options with trade-offs, costs, risks, and the assumptions behind each. Do not pick a winner unless I ask.
- What would a domain expert in sales, legal, operations, finance, data science, or regulation criticise first? Use explicit criteria rather than pretending to be that expert.
- Rewrite this for a CEO, customer, engineer, employee, regulator, or board without removing the hard parts.
- What is the cheapest safe version of this idea? What is the smallest reversible pilot?
- What deterministic tool is safer than you for each part of this task?
- Return the output in a structured schema with fields for evidence, assumptions, uncertainty, risks, action owner, and next step.
- Show calculations, code, transformations, and decision rules so another person can reproduce the result.
- Create a checklist I should run before sending this output to a client, manager, customer, or regulator.
- What constraints, examples, counterexamples, or format should I add next time to improve the result?
- Challenge my framing: am I asking the wrong question?

14 Reproducibility and organisational memory

When AI becomes part of a recurring process, the organisation needs to remember not only the answer, but how the answer was produced and corrected.

- Which model, version, prompt, tools, retrieval sources, data snapshot, and permissions produced this result?
- Can another person reproduce the answer from the same inputs?

- What changed between this answer and the previous version?
- How are corrections, overrides, rejected recommendations, and incidents recorded?
- Which documents or facts are the authoritative source of truth?
- How are stale, incorrect, duplicated, or harmful memories removed or corrected?
- What knowledge should be sedimented into a shared organisational asset rather than remaining in a private chat?
- Which recurring exceptions reveal that the workflow, policy, data, or model needs redesign?

15 Better prompting: ask the AI how to use it

These prompts improve the interaction itself:

- Turn this messy brief into a structured prompt I can reuse.
- What context is missing from this prompt, and which missing item has the highest expected impact?
- Ask me at most five high-value questions. If I do not answer, continue with explicit assumptions.
- Show me a checklist I should run before I send your output to a client.
- Suggest examples and counterexamples that would make the task unambiguous.
- What constraints, data, tools, or output format should I add next time?
- What part of this task should be delegated to search, code, a database, a calculator, or a human expert?
- What would a good answer refuse to claim?
- Which parts of your answer should I independently verify first?
- Challenge my framing: am I asking the wrong question?

16 Reusable prompt templates

16.1 The default business prompt

Here is the business context: [CONTEXT].
The decision I need to make is: [DECISION].
The people affected are: [STAKEHOLDERS].
The available data and tools are: [DATA AND TOOLS].
The constraints are: [CONSTRAINTS].
The deadline and audience are: [DEADLINE AND AUDIENCE].
The cost of being wrong is: [CONSEQUENCE].

Do not invent facts or sources. Separate my inputs, verified evidence, assumptions, inferences, and unknowns. Challenge my framing, identify missing information, compare alternatives, state the main risks, define what a human must approve, and propose the cheapest test that could disprove the recommendation. If the evidence is insufficient, say so and recommend whether to proceed, pilot, escalate, or stop.

16.2 The anti-sycophancy prompt

Do not assume that my framing is correct. First restate the problem neutrally and identify the assumptions embedded in my request. Give me at least three plausible interpretations or options. Separate facts, evidence, assumptions, inferences, and unknowns. Mark missing information with [MISSING]. Do not invent details to make the answer complete. Explain what evidence would change your conclusion and where a human expert must decide.

16.3 The pilot-design prompt

Design a small, reversible pilot for [USE CASE]. Define the current baseline, the decision that should change, the users, the AI and human responsibilities, the data required, the main failure modes, the acceptance criteria, the monitoring metrics, the fully loaded cost, and the stopping rule. Include normal, edge, adversarial, and failure cases. State what result would make us scale, redesign, pause, or stop.

16.4 The agent-permissions prompt

For an AI agent performing [TASK], specify the minimum permissions it needs, what it may read, write, send, approve, purchase, delete, or publish, and what is permanently off-limits. Add approval gates for irreversible actions, logging requirements, prompt-injection defences, a manual fallback, an incident process, and a safe rollback plan. Assume that retrieved documents and external messages may contain untrusted instructions.

16.5 The source-grounded prompt

Answer using only the supplied documents and sources dated between [DATE RANGE], relevant to [JURISDICTION]. For each important claim, provide the source, date, and a short explanation of how directly it supports the claim. If the sources disagree, show the disagreement. If a claim is not supported, mark it [SOURCE NEEDED]; do not fill the gap from general knowledge.

17 Before acting: the final gate

Before using an AI output in the real world, ask:

1. What decision is this output informing?
2. What evidence supports the important claims?
3. What remains unknown or uncertain?
4. What could go wrong, and who would bear the cost?
5. Has a suitable human reviewed it?
6. Is the output being used within its authorised purpose and risk level?
7. Can the action be reversed if the output is wrong?
8. Have I checked personal data, confidential information, IP, legal, and security issues?
9. What baseline and metric will tell us whether this helped?
10. When do we stop, redesign, or escalate?

Final takeaway. The best AI question is not simply “What is the answer?” It is: “What decision are we making, what do we know, what do we not know, what alternatives remain possible, how can we test this cheaply, and who is accountable for acting on it?”